

Sublinear-Space Approximations of Vertex Centrality in Evolving Graphs

Benjamin W. Priest

Thayer School of Engineering
Dartmouth College


benjamin.w.priest.th@dartmouth

December 6, 2018

- 1 Introduction
- 2 Background
- 3 Pseudo-Asynchronous Communication Schemes for Vertex-Centric Distributed Algorithms
- 4 DEGREE SKETCH and Local Triangle Count Heavy Hitters
- 5 Sublinear κ -Path Centrality


Motivation

- Many modern computing problem focus on complex relational data
- Data are phrased as large graphs
 - e.g. the Internet, communication networks, transportation systems, protein networks, epidemiological models, social networks
- Often want to identify which vertices are “important”
 - Robust to changes?
 - Sublinear memory?




Overcoming Data Scale: Data Streaming

- Traditional RAM algorithms scale poorly
 - Awkward to store data in memory
 - Superlinear scaling unacceptable
- Data stream model to the rescue!
 - Sequential data access
 - Sublinear memory
 - Linear amortized time
 - Constrained number of passes
 - Monte Carlo Approximations




Overcoming Data Scale: Distributed Data Streaming

- Distributed memory model a staple of HPC
 - Divide computation across many processors
 - Communication an important resource
 - Immense scaling of exact algorithms
- Why not distributed data streams!
 - Sketch data structures afford stream composition
 - Works nicely with vertex-centric algorithms
 - Even greater scaling
 - Linear communication



Centrality Indices

- Assign scores to vertices
 - Higher score \rightarrow more important
 - Depends on graph structure
 - Different indices in different domains
- Scores are not informative
 - Usually want top k vertices
- Relative order-preserving approximation is acceptable



The Problem

- Memory overhead
- Computational Overhead
- Communication Overhead
- Wasted effort
 - Generally only need top elements vis-à-vis a centrality index

Our Solution


- Sketch data structures
 - Utilize composable streaming summaries of vertex-local information
- Distributed memory
 - Partition graph and distribute sketches
 - Polyloglinear computation, memory, and communication

Graph Primitives

Assume throughout that $\mathcal{G} = (\mathcal{V}, \mathcal{E}, \mathbf{w})$, where $|\mathcal{V}| = n$ and $|\mathcal{E}| = m$

- \mathbf{w}_e is the weight of edge e if $e \in \mathcal{E}$ and zero otherwise
- \mathcal{G} has adjacency matrix $A \in \mathbb{R}^{n \times n}$ so that $A_{x,y} = \mathbf{w}_{xy}$ for $xy \in \mathcal{E}$
- \mathcal{G} has vertex-edge incidence matrix $B \in \mathbb{R}^{\binom{n}{2} \times n}$ so that

$$B_{xy,z} = \begin{cases} \mathbf{w}_{xy} & \text{if } x = z \\ -\mathbf{w}_{xy} & \text{if } y = z \\ 0 & \text{else.} \end{cases}$$



$$B = \begin{matrix} & \begin{matrix} a & b & c & d \end{matrix} \\ \begin{matrix} ab \\ ac \\ ad \\ bc \\ bd \\ cd \end{matrix} & \begin{pmatrix} 1 & -1 & 0 & 0 \\ 1 & 0 & -1 & 0 \\ 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 \\ 0 & 1 & 0 & -1 \\ 0 & 0 & 1 & -1 \end{pmatrix} \end{matrix}$$

Streaming Background

- A *stream* σ accumulating $\mathbf{f} \in \mathbb{R}^n$ is a list of rank-1 updates
 - An update (i, c) means $M \leftarrow \mathbf{f} + c * \mathbf{e}_i$
 - A *cash register* stream enforces $c > 0$ for all updates
 - A *turnstile* stream allows negative updates
 - A *strict turnstile* stream allows negatives but enforces $\mathbf{f} \in \mathbb{R}_{\geq 0}^n$
- An algorithm accumulating a data structure \mathcal{S} and is said to be...
 - *streaming* if \mathcal{S} uses $O(\log n)$ memory
 - *semi-streaming* if \mathcal{S} uses $O(n \text{ polylog } n)$ ¹ memory
- Want to minimize the number of passes over σ
 - 1 pass ideal
 - Constant or logarithmic passes sometimes acceptable

¹sometimes $O(n^{1+\alpha})$ for $\alpha \in (0, 1/2]$

Sketching

Definition (Sketch)

A *Sketch* is a streaming data structure \mathcal{S} that admits a merge operator \oplus . If \circ is the stream concatenation operator, then for any streams σ_1 and σ_2 ,

$$\mathcal{S}(\sigma_1) \oplus \mathcal{S}(\sigma_2) = \mathcal{S}(\sigma_1 \circ \sigma_2).$$

Definition (Linear Sketch)

A *Linear Sketch* \mathcal{S} is a linear projection of \mathcal{U} to a lower dimension. For any streaming frequency vectors \mathbf{f}_1 and \mathbf{f}_2 and scalars a and b ,

$$a\mathcal{S}(\mathbf{f}_1) + b\mathcal{S}(\mathbf{f}_2) = \mathcal{S}(a\mathbf{f}_1 + b\mathbf{f}_2).$$

Sketches are useful for stream summarization when
comparisons between streams are important

Summary of Results: Serial Algorithms

Degree Centrality

$$\mathcal{C}^{\text{DEG}}(x) = |\{(u, v) \in E \mid x \in \{u, v\}\}| = \|A_{x,:}\|_1 = \|A_{:,x}\|_1$$

- Naïve online $O(n)$ -space and -time algorithm exists

Summary of Results: Serial Algorithms

Degree Centrality

We show $\tilde{O}(1)$ -space distributable streaming algorithms

- Naïve online $O(n)$ -space and -time algorithm exists

Summary of Results: Serial Algorithms

Degree Centrality

We show $\tilde{O}(1)$ -space distributable streaming algorithms

- Naïve online $O(n)$ -space and -time algorithm exists

Closeness Centrality

$$c^{\text{CLOSE}}(x) = \frac{1}{\sum_{y \in V} d(x, y)}$$

- Online exact $O(n^2)$ -space $O(nm)$ -time algorithm [WC14]
- Batch Approximate $O(n^2)$ -space and almost-linear time algorithm [CDPW14]

Summary of Results: Serial Algorithms

Degree Centrality

We show $\tilde{O}(1)$ -space distributable streaming algorithms

- Naïve online $O(n)$ -space and -time algorithm exists

Closeness Centrality

$$C^{\text{CLOSE}}(x) = \frac{1}{\sum_y d(x, y)}$$

We show constant-pass semi-streaming algorithm

- Online
- Batch Approximate $O(n^2)$ -space and almost-linear time algorithm [CDPW14]

Summary of Results: Distributed Streaming Algorithms

Triangle Count Centrality

$$\mathcal{C}^{\text{TRI}}(x) = |\{yz \in \mathcal{E} \mid xy, yz, xz \in \mathcal{E}\}| \quad (\text{vertex-local})$$

$$\mathcal{C}^{\text{TRI}}(xy) = |\{z \in \mathcal{E} \mid xy, yz, xz \in \mathcal{E}\}| \quad (\text{edge-local})$$

- Exact $O(m)$ -space, $O(m^{\frac{3}{2}})$ serial and distributed algorithms [AKM13]
- Streaming sampling sublinear-space algorithms [LK15, SERU17]
 - Including distributed generalizations [SHL⁺18, SLO⁺18, PPS18]

Summary of Results: Distributed Streaming Algorithms

Triangle Count Centrality

$$\mathcal{C}^{\text{TRI}}(x) = |\{yz \in \mathcal{E} \mid xy, yz, xz \in \mathcal{E}\}| \quad (\text{vertex-local})$$

We show 2-pass, semi-streaming, distributed sketch-based query algorithms for estimating heavy hitters

- Ex [KM13]
- Streaming sampling sublinear-space algorithms [LK15, SERU17]
 - Including distributed generalizations [SHL⁺18, SLO⁺18, PPS18]

Summary of Results: Distributed Streaming Algorithms

Triangle Count Centrality

$$C^{\text{TRI}}(x) = |\{yz \in \mathcal{E} \mid xy, yz, xz \in \mathcal{E}\}| \quad (\text{vertex-local})$$

We show 2-pass, semi-streaming, distributed sketch-based query algorithms for estimating heavy hitters

- Example: ϵ -approximation of C^{TRI} in $\tilde{O}(m)$ space [KM13]
- Streaming sampling sublinear-space algorithms [LK15, SERU17]
 - Including distributed generalizations [SHL⁺18, SLO⁺18, PPS18]

κ -Path Centrality

$$C^{\kappa}(x) = \Pr_{p: |p| \leq \kappa} [x \in p \wedge p \text{ a simple path}]$$

- $O(m)$ -space $O(n^{1+\alpha} \log^2 n)$ -time approximation algorithm [WC14]
- Empirical proxy for betweenness centrality heavy hitters
 - Online exact and approximate $O(n^2)$ - and $O(m)$ -space algorithms exist [GMB12, WC14, KMB15, BMS14]

Summary of Results: Distributed Streaming Algorithms

Triangle Count Centrality

$$C^{\text{TRI}}(x) = |\{yz \in \mathcal{E} \mid xy, yz, xz \in \mathcal{E}\}| \quad (\text{vertex-local})$$

We show 2-pass, semi-streaming, distributed sketch-based query algorithms for estimating heavy hitters

- Estimating triangle count centrality [KM13]
- Streaming sampling sublinear-space algorithms [LK15, SERU17]
 - Including distributed generalizations [SHL⁺18, SLO⁺18, PPS18]

κ -Path Centrality

$$C^{\kappa}(x) = \Pr_{p \leftarrow \mathcal{P}} [x \in p \wedge p \text{ a simple path}]$$

We show distributed sublinear vertex-centric sampling algorithm

- $O(n)$ space algorithm [C14]
- Empirical proxy for betweenness centrality heavy hitters
 - Online exact and approximate $O(n^2)$ - and $O(m)$ -space algorithms exist [GMB12, WC14, KMB15, BMS14]

Motivation: Vertex-Centric Algorithms

The Problem:

- Most distributed graph algorithms are vertex-centric
 - Partition local vertex information across processors
 - Processors communicate as in rounds $[MAB^{+10}]$
- Scale-free graphs common in applications
 - Exhibit very high degree vertices
 - Cause computation, communication, and memory “hotspots”
 - Synchronous communication moves at the speed of the slowest processor

Existing Solutions:

- Asynchronous Communication
 - Processors communicate point-to-point as needed
 - Increased implementation complexity
- Vertex delegation [PGA14]
 - Subpartition high degree vertices between processors
 - Adds communication overhead

Approach: Pseudo-Asynchronous Communication Protocol

The Idea

Allow processors to drop out of communication exchanges when finished

- Partition processor set \mathcal{P} into *local* and *remote* exchanges
 - Takes advantage of hybrid distributed memory
- $P \in \mathcal{P}$ maintains send buffer $\mathcal{S}[P]$ and receive buffer $\mathcal{R}[P]$
 - Begin forwarding messages when $\mathcal{S}[P]$ reaches threshold
 - Drop out of exchange when all other processors in exchanges stop updating $\mathcal{R}[P]$
- Three protocols:
 - Node Local
 - Node Remote
 - Node Local Node Remote (NLNR)



Node Local and Node Remote

Node Local

- 1 Send messages to local core matching destination local offset
- 2 Send messages to remote core matching destination node offset
- 3 Good for mostly point-to-point messages



Node Remote

- 1 Send messages to remote core matching destination node offset
- 2 Send messages to local core matching destination local offset
- 3 Good for large numbers of broadcasts



Node Local Node Remote

- Further partition processors by *layers*
 - A layer is a collection of nodes equal to the number of cores per node
- ❶ Send messages to local core matching destination layer offset
- ❷ Send messages to remote core matching destination node offset
- ❸ Send messages to local core matching destination local offset
- ❹ Best for extreme scale where many layers exist

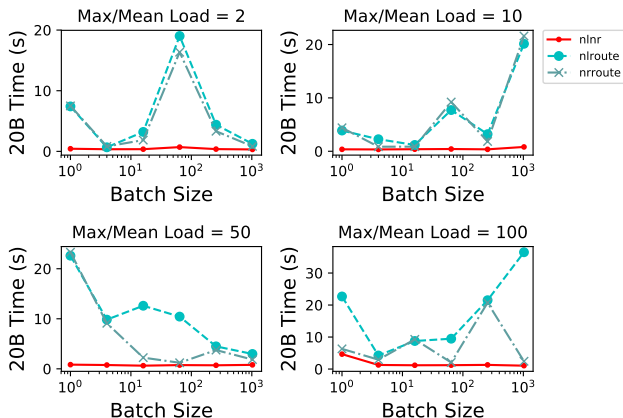


Validation of Claims

Experiment

- 20B message exchange
 - Destination sampled from Pareto distribution
 - Batch size is maximum $|S[P]|$

Experiment on N=512 nodes with C=32 cores



NLNR exhibits best scaling, others more useful with fewer nodes

YGM C++/MPI Library

- Authored by myself, Trevor Steil (UMN), and Roger Pearce (LLNL)
- Simple API for handling pseudo-asynchronous communication
 - Clients need only specify receive behavior
- Supports message serialization for arbitrary, variable-length messages
- Supports LLNL Projects
 - HAVOQgt
 - graph500 scale leader
 - others?
- Useful for not just vertex-centric algorithms, but any algorithm with asymmetric computational and communication load

YGM to be open sourced

Motivation: Local Triangle Counting

The Problem:

- Local triangle counting a common big data analytic
 - Exact computation expensive $O\left(m^{\frac{3}{2}}\right)!$
- Recall

$$\mathcal{C}^{\text{TRI}}(x) = |\{yz \in \mathcal{E} \mid xy, yz, xz \in \mathcal{E}\}| \quad (\text{vertex-local})$$

$$\mathcal{C}^{\text{TRI}}(xy) = |\{z \in \mathcal{E} \mid xy, yz, xz \in \mathcal{E}\}| \quad (\text{edge-local})$$



Existing Solutions:

- Many exact distributed algorithms [AKM13, Pea17]
- Many approximate streaming algorithms via sampling [LK15, SERU17]
- ... and some utilizing both models [SHL⁺18, SLO⁺18]

Approach: Sublinear Intersection Method

Idea: Intersection method, but using cardinality sketches

- Cardinality sketches summarize set size
- Support union operation, and some support limited intersection operation
 - High variance if intersection is small
 - Likely best performance on heavy hitters
- Affords edge- and vertex-local triangle count estimation
- Outputs only reliable if *triangle density* is nontrivial
 - Triangle density = $\frac{\# \text{ triangles}}{\# \text{ possible triangles}}$



HYPERLOGLOG Cardinality Sketches

HLL cardinality sketches

Maintain $r = 2^p$ 6-bit registers M and a 64-bit hash function h

- Insert x : let $i = \langle x_1, \dots, x_p \rangle$ and $w = \langle x_{p+1}, \dots, x_{64} \rangle$
- $\rho(w)$ = initial zero bits of w plus 1
- $M_i = \max\{M_i, \rho(w)\}$
- Estimator derives from harmonic mean of M

HYPERLOGLOG Cardinality Sketches

HLL cardinality sketches

Maintain $r = 2^p$ 6-bit registers M and a 64-bit hash function h

- Insert x : Outputs \tilde{C} such that for cardinality C ,
w.h.p. $|C - \tilde{C}| \leq \frac{1.04}{\sqrt{m}} C$ [FFGM07]
- $\rho(w) = i$
- $M_i = \max\{M_i, \rho(w)\}$
- Estimator derives from harmonic mean of M

HYPERLOGLOG Cardinality Sketches

HLL cardinality sketches

Maintain $r = 2^p$ 6-bit registers M and a 64-bit hash function h

- Insert x : Outputs \tilde{C} such that for cardinality C ,
w.h.p. $|C - \tilde{C}| \leq \frac{1.04}{\sqrt{m}} C$ [FFGM07]
- $\rho(w) = i$
- $M_i = \max\{M_i, \rho(w)\}$
- Estimator derives from harmonic mean of M

Useful results

- Native intersection operator (elementwise maximum)
- Various improved harmonic [HNH13, QKT16] and maximum likelihood estimators [XZC17, Lan17, Ert17]
- Sparsification for low cardinality sets [HNH13]
- Compression to 4 and 3 bit registers [XZC17]
- Intersection estimators [Tin16, CKY17, Ert17]

DEGREE SKETCH and Triangle Counting

Assume a partition $f : \mathcal{V} \rightarrow \mathcal{P}$, and let $\mathcal{V}_P = \{v \in \mathcal{V} \mid f(v) = P\}$

- Distribute DEGREE SKETCH \mathcal{D} across \mathcal{P}
 - $\mathcal{D}[v]$ holds a HLL for adjacency set of $v \in \mathcal{V}$
 - P holds $\mathcal{D}[v]$ for $v \in \mathcal{V}_P$
- Accumulate \mathcal{D} in one pass over σ
 - Assume $P \in \mathcal{P}$ gets substream σ_P
 - P sends $xy \in \sigma_P$ to $f(x)$ and $f(y)$
 - When P gets $xy : x \in \mathcal{V}_P$, insert y into $\mathcal{D}[x]$
 - $\mathcal{D}[x]$ starts sparse and eventually saturates
- \mathcal{D} can be queried after estimation, e.g.
 - Estimate $\tilde{\mathcal{C}}^{\text{DEG}}(v) = \text{ESTIMATE}(\mathcal{D}[v])$
 - Estimate $\tilde{\mathcal{C}}^{\text{TRI}}(uv) = \mathcal{D}[u] \tilde{\cap} \mathcal{D}[v]$
 - Involves communication if $f(u) \neq f(v)$
 - Estimate $\tilde{\mathcal{C}}^{\text{TRI}}(v) = \frac{\sum_{uv \in \mathcal{E}} \tilde{\mathcal{C}}^{\text{TRI}}(uv)}{2}$
 - Requires second pass in general

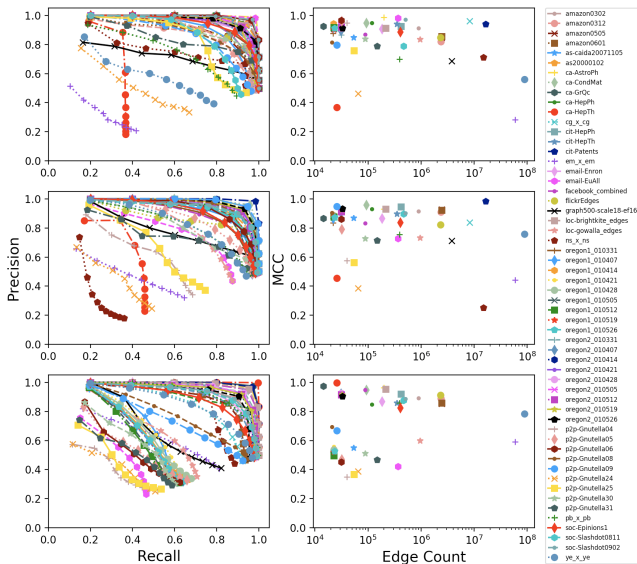
$\tilde{O}(m)$ time and communication and $\tilde{O}(\varepsilon^{-2}n)$ space!

Edge-Local Triangle Count Heavy Hitters

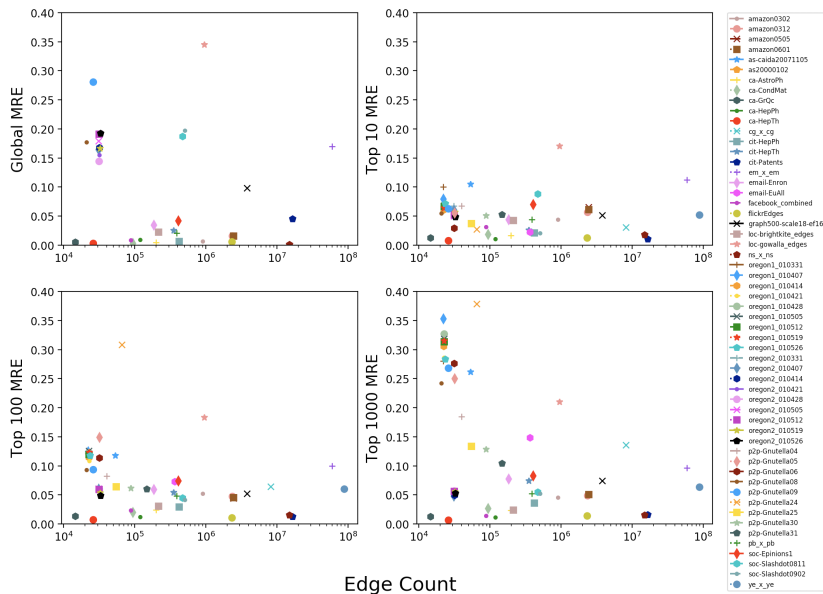
Algorithm 1 Edge-Local Triangle Count Heavy Hitters

```
1: Accumulate  $\mathcal{D}$  in distributed pass over  $\sigma$ 
2:  $H_k \leftarrow$  empty  $k$ -heap
3:  $T \leftarrow 0$ 
4: parallel for  $xy \in \sigma_P$  do           // second pass
5:   Send  $(E, xy)$  to  $f(x)$  and  $(E, yx)$   $f(y)$ 
6:   for  $(E, xy) \in \mathcal{R}[P]$  do
7:     Send  $(S, xy, \mathcal{D}[x])$  to  $f(y)$ 
8:     for  $(S, xy, \mathcal{D}[x]) \in \mathcal{R}[P]$  do
9:       Insert  $(xy, \mathcal{D}[x] \tilde{\cap} \mathcal{D}[y])$  into  $H_k$ 
10:       $T \leftarrow T + \mathcal{D}[x] \tilde{\cap} \mathcal{D}[y]$ 
11:  $T \leftarrow T/2$ 
12: Global accounting of  $T, H_k$ 
13: return  $H_k$ 
```





Validation of Claims: Precision and Recall



Validation of Claims: Relative Error



Validation of Claims: Good vs Bad Triangle Density



Triangle Count Heavy Hitters

DEGREEskETCH C++/MPI Library

- Authored by myself
- Utilizes YGM for communication
- Accumulation and query API for DEGREEskETCH
- Supports sparse and compressed registers
- Implementations for edge- and vertex-local triangle count heavy hitter estimation
- Supports more exotic queries

DEGREEskETCH to be open sourced

Motivation: Betweenness Centrality Heavy Hitters

The Problem:

- Computing Betweenness centrality exactly amounts to computing `ALLSOURCESALLSHORTESTPATHS`
 - Expensive $O(mn)$!

Existing Solutions:

- Approximate via a logarithmic number of `SINGLESOURCEALLSHORTESTPATHS` [GMB12, BMS14, Yos14, KMB15, RK16]
 - Difficult to distribute
 - Unclear if possible in $o(m)$ memory

Approach: Sublinearize κ -Path Centrality

Idea: “Come at the problem sideways”

- High κ -path centrality empirically correlates with high betweenness centrality [KAS⁺13]
- Algorithm amounts to sampling random simple paths
 - Use ℓ_p sampling sketches to sublinearize
- Sublinear approximation of κ -path centrality \rightarrow empirical recovery of high betweenness centrality vertices?

κ -path centrality

$$PC(x, \kappa) = \Pr_{p: |p| \leq \kappa} [x \in p \wedge p \text{ a simple path}]$$

“simple path” = non-self-intersecting path

ℓ_p Sampling Sketches

ℓ_p sampling sketches

Sample from frequency vector \mathbf{f} with probability relative to ℓ_p norm

- Sample $t_i \sim_R (0, 1) \forall i \in [n]$
- Rescale updates to \mathbf{f}_i by $1/t_i^{1/p}$
- Accumulate TUG-OF-WAR, COUNTSKETCH, and ℓ_p norm sketches
- Use sketches to output COUNTSKETCH argmax or FAIL

ℓ_p Sampling Sketches

ℓ_p sampling sketches

Sample from frequency vector \mathbf{f} with probability relative to ℓ_p norm

- Sample $t_i \sim_R (0, \dots, \mathbf{f}_i)$
- Rescale updates
- Accumulate TUG-OF-WAR, COUNTSKETCH, and ℓ_p norm sketches
- Use sketches to output COUNTSKETCH argmax or FAIL

Blah blah etc etc
It is complex

ℓ_p Sampling Sketches

ℓ_p sampling sketches

Sample from frequency vector \mathbf{f} with probability relative to ℓ_p norm

- Sample t_i Outputs (i, P) w.p. $1 - \delta$, where $i \in [n]$ is sampled w.p. $P = (1 \pm \varepsilon) \frac{|v_i|^p}{\|\mathbf{v}\|_p^p}$ [MW10]
- Rescale u_i
- Accumulate TUG-OF-WAR, COUNTSKETCH, and ℓ_p norm sketches
- Use sketches to output COUNTSKETCH argmax or FAIL

Useful results

[JST11, Vu18]

- ℓ_0 sketch requires $\tilde{O}(\log(1/\delta))$ memory and update time
 - Useful for unweighted random hops
- ℓ_1 sketch requires $\tilde{O}(\varepsilon^{-1} \log(1/\delta))$ memory and $\tilde{O}(\log(1/\delta))$ update time
 - Useful for weighted random hops
- s parallel ℓ_p sketches can be accumulated in time independent of s

ℓ_p Sampling Graph Sparsification

- Exploit sketch linearity
 - Sample ℓ_0 sampling sketch matrices S_1, \dots, S_t , each of which will sketch every column of B
 - $S_1(B_{:,x})$ returns a sampled neighbor of x , say y
 - $S_2(B_{:,x}) + S_2(B_{:,y}) = S_2(B_{:,x} + B_{:,y})$ returns a sampled neighbor of the supervertex $(x + y)$
 - et cetera
- This method can solve several problems [AGM12a, AGM12b]:
 - $O(n \text{ polylog } n)$ to decide connectivity, k -connectivity, bipartiteness, and to approximate the weight of the MST
 - Multipass $\tilde{O}(n^{1+1/\alpha})$ to compute sparsifiers, the exact MST, α -**spanners**, and approximate the maximum weight matching

We will use similar methods to sample random walks and random simple paths in distributed algorithms

Distributed Accumulation ℓ_p Sampling Sketches

- $P \in \mathcal{P}$ accumulates adjacency set $\mathcal{A}[v]$ for each $v \in \mathcal{V}_P$
 - When $\mathcal{A}[v]$ too large, replace it with s ℓ_0 sampling sketches
 - Write current state and all subsequent updates to disk memory
- Queries to $\mathcal{A}[v]$ return a sampled neighbor of v
 - If $\mathcal{A}[v]$ is a set of sketches, one is consumed
 - If FAIL, repeat
 - Once $\mathcal{A}[v]$ sketches are exhausted, P takes another pass over v 's substream in disk memory

Avoids need to subpartition vertices across multiple processors, effectively exchanging communication time for I/O time

Sublinear Random Walk and Simple Path Sampling

Random Walk Simulation

- Sample t vertices $\{v_{1,1}, \dots, v_{t,1}\}$ and:
 - Sample $v_{i,j+1}$ from $\mathcal{A}[v_{i,j}]$
 - Communicate $(v_{i,1}, \dots, v_{i,j+1})$ to $f(v_{i,j+1})$

Random Simple Path Simulation

- Similar to random walks, except:
 - Do not accumulate sketches ahead of time
 - Sample $v_{i,j+1}$ from $\mathcal{A}[v_{i,j}] \setminus \{v_{i,1}, \dots, v_{i,j-1}\}$
 - If $\mathcal{A}[v_{i,j}]$ is not in memory, accumulate a sketch ignoring edges to any of $\{v_{i,1}, \dots, v_{i,j-1}\}$
 - Communicate $(v_{i,1}, \dots, v_{i,j+1})$ to $f(v_{i,j+1})$

Sublinear distributed storage of graph by sketching high degree vertices

Sublinear κ -Path Centrality

κ -Path Centrality Approximation Algorithm ([KAS⁺13]):

- ① Simulate $T = 2\kappa^2 n^{1-2\alpha} \ln n$ ($\leq \kappa$)-length simple paths over \mathcal{G}
 - maintain $\text{count}[x]$ for each $x \in \mathcal{V}$
- ② $\tilde{\mathcal{C}}^\kappa(x) \leftarrow \frac{\text{count}[x]}{2\kappa n^{-2\alpha} \ln n}$

- Given $\alpha \in [-1/2, 1/2]$, for each $x \in \mathcal{V}$, $|\tilde{\mathcal{C}}^\kappa(x) - \mathcal{C}^\kappa(x)| \leq n^{1/2+\alpha}$ w.h.p.
- Easy to distribute in vertex-centric model

Sublinear κ -Path Centrality

Algorithm 2 Sublinear κ -Path Centrality

```
1: for  $i \in \{1, \dots, T\}$  do
2:    $p_i \leftarrow$  empty path
3:    $p_{i,1} \leftarrow$  uniform sample from  $\mathcal{V}$ 
4:    $l_i \leftarrow$  uniform sample from  $\{1, 2, \dots, \kappa\}$ 
5:   for  $x \in \mathcal{V}$  do  $c_x \leftarrow 0$ 
6:   parallel for  $j \in \{1, 2, \dots, \kappa - 1\}$  do
7:     parallel for  $i \in \{1, 2, \dots, T\}$  do
8:       if  $j < l_i$  then
9:          $p_{i,j+1} \leftarrow$  sample from  $\mathcal{A}[p_{i,j}]$ 
10:        if  $p_{i,j+1} = \emptyset$  then discard
11:      else if  $j = l_i$  then
12:         $c_{p_{i,k}} \leftarrow c_{p_{i,k}} + 1$  for  $k \in \{1, \dots, j\}$ 
13: return  $c_x / 2\kappa n^{-2\alpha} \ln n$  for  $x \in V$ 
```

Summary of Results

- The Goal: distributed sublinear approximations of centrality indices
- Engineering Results
 - YGM: Pseudo-Asynchronous Communication Handler
- Algorithmic Results
 - A streaming degree centrality approximation and heavy hitter recovery algorithms
 - A $O(1)$ -pass semi-streaming closeness centrality approximation algorithm
 - 2-pass distributed semi-streaming edge- and vertex-local triangle count heavy hitter estimation algorithms using `DEGREE SKETCH`
 - Distributed sublinear semi-streaming random walk and random simple path sampling algorithms
 - Distributed sublinear semi-streaming κ -path centrality estimation algorithm
- Future Work
 - Applications for `DEGREE SKETCH`
 - Sublinear random walk and random simple path implementation

Questions?

